

Reinforcement learning

- Reinforcement Learning is the agent must sense the environment, learns to behave (act) in a environment by performing actions (reinforcement) and seeing the results.
- Task
 - Learn how to behave successfully to achieve a goal while interacting with an external environment.
 - The goal of the agent is to **learn** an **action policy** that maximizes the total reward it will receive from any starting state.
- Examples
 - **Game playing:** player knows whether it win or lose, but not know how to move at each step

Reinforcement Learning Process

- RL contains two primary components:
 1. Agent (A) – RL algorithm that learns from trial and error
 2. Environment – World Space in which the agent moves (interact and take action)
- State (S) – Current situation returned by the environment
- Reward (R) – An immediate return from the environment to appraise the last action
- Policy (π) – Agent uses this approach to decide the next action based on the current state
- Value (V) – Expected long-term return with discount. Oppose to the short-term reward (R)
- Action-Value (Q) – Similar to value except it contains an additional parameter, the current action (A)



RL Approaches


- Two approaches
 - *Model based approach RL:*
 - learn the model, and use it to derive the optimal policy.
e.g Adaptive dynamic learning(ADP) approach
 - *Model free approach RL:*
 - derive the optimal policy without learning the model.
e.g LMS and Temporal difference approach
- Passive learning
 - The agent simply watches the world during transition and tries to learn the utilities in various states
- Active learning
 - The agent not simply watches, but also acts on the environment

Reinforcement learning model

- Each percept(e) is enough to determine the State(the state is accessible)
- **Agent's task:** Find a optimal policy by mapping states of environment to actions of the agent, that maximize long-run measure of the reward (reinforcement)
- It can be modeled as Markov Decision Process (MDP) model.
 - Markov decision process (**MDP**) is a a mathematical framework for **modeling** decision making i.e mapping a solution in reinforcement learning.

MDP model

- MDP model $\langle S, T, A, R \rangle$




- S – set of states
- A – set of actions
- Transition Function:** $T(s, a, s') = P(s'|s, a)$ – the probability of transition from s to s' given action a
 $T(s, a) \rightarrow s'$
- Reward Function:** $r(s, a) \rightarrow r$ the expected reward for taking action a in state s

$$R(s, a) = \sum_{s'} P(s'|s, a) r(s, a, s')$$

$$R(s, a) = \sum_{s'} T(s, a, s') r(s, a, s')$$

MDP - Example I

- Consider the graph, and find the shortest path from a node S to a goal node G.
- Set of states $\{S, T, U, V\}$
- **Action** – Traversal from one state to another state
- **Reward** - Traversing an edge provides "length edge" in dollars.
- **Policy** – Path considered to reach the destination $\{S \rightarrow T \rightarrow V\}$



Q - Learning

- **Q-Learning** is a value-based **reinforcement learning** algorithm uses Q-values (action values) to iteratively improve the behavior of the learning agent.
- Goal is to maximize the Q value to find the optimal action-selection policy.
- The Q table helps to find the best action for each state and maximize the expected reward.
- **Q-Values / Action-Values:** Q-values are defined for states and actions.
- $Q(s, a)$ denotes an estimation of the action a at the state s .
- This estimation of $Q(s, a)$ will be iteratively computed using the **TD-Update rule**.
- **Reward:** At every transition, the agent observes a reward for every action from the environment, and then transits to another state.
- **Episode:** If at any point of time the agent ends up in one of the terminating states i.e. there are no further transition possible is called **completion of an episode**.

Q – Learning Algorithm

- Set the gamma parameter
- Set environment rewards in matrix R
- Initialize matrix Q as Zero
 - Select random initial (source) state
 - Set initial state $s = \text{current state}$
 - Select one action a among all possible actions using exploratory policy
 - Take this possible action a , going to the *next state* s' .
 - Observe reward r
 - Get maximum Q value to go to next state based on all possible actions
- Compute:
 - $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$
- Repeat the above steps until reach the goal state i.e current state = goal state

Understanding the Q – Learning: **Prepare matrix Q**

- Matrix Q is the memory of the agent in which learned information from experience is stored.
- Row denotes the current state of the agent
- Column denotes the possible actions leading to the next state


Compute Q matrix:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$$

- Gamma is discounting factor for future rewards. Its range is 0 to 1. i.e. $0 < \text{Gamma} < 1$.
- Future rewards are less valuable than current rewards so they must be discounted.
- If Gamma is closer to 0, the agent will tend to consider only the immediate rewards.
- If Gamma is closer to 1, the agent will tend to consider only future rewards with higher edge weights.



Understanding the Q – Learning

- Building Environment contains 5 rooms that are connected with doors.
- Each room is numbered from 0 to 4. The building outside is numbered as 5.
- Doors from room 1 and 4 leads to the building outside 5.
- **Problem:** Agent can place at any one of the rooms (0, 1, 2, 3, 4). Agent's goal is to reach the building outside (room 5).




Understanding the Q – Learning

- Represent the room in the graph.
- Room number is the state and door is the edge.




Understanding the Q – Learning

- Assign the Reward value to each door.
- The doors lead immediately to target is assigned an instant reward of 100.
- Other doors not directly connected to the target room have **zero** reward.
- For example, doors are two-way (0 leads to 4, and 4 leads back to 0), two edges are assigned to each room.
- Each edge contains an instant reward value



Understanding the Q – Learning

- Let consider agent starts from state s (Room) 2.
- Agent's movement from one state to another state is action a .
- Agent is traversing from state 2 to state 5 (Target).
 - Initial state = current state i.e. state 2
 - Transition State 2 \rightarrow State 3
 - Transition State 3 \rightarrow State (2, 1, 4)
 - Transition State 4 \rightarrow State 5




Understanding the Q – Learning

Prepare rewards table R (matrix)

- -1 denotes the no edge between the states
- 0 represents the indirect edge to the target

$$R = \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix}$$



Example: Q – Learning

Matrix Q :

- Set the Gamma value = 0.8
- Initialize the matrix Q to zero matrix

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

Example: Q – Learning


Matrix Q :

- Set the Gamma value = 0.8
- Initialize the matrix Q to zero matrix

$$Q = \begin{matrix} & \begin{matrix> \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix> \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix> \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix> \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

- For next episode, next state 1 becomes current state
- Repeat the inner loop due to 1 is not target state
- From State 1, either can go to 3 or 5.
- Let's choose state 5.
- Compute max Q value to go to **next state** based on all possible actions.



- $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \max[Q(\text{next state}, \text{all actions})]$
- $Q(1,5) = R(1,5) + 0.8 * \max[Q(5,1), Q(5,4), Q(5,5)]$
 $= 100 + 0.8 * \max[0, 0, 0] = 100 + 0 = \mathbf{100}$
- Q remains the same due to $Q(1,5)$ is already fed into the agent. Stop process

0

	0	1	2	3	4	5		State	0	1	2	3	4	5
$R =$	0	-1	-1	-1	0	-1	$Q =$	0	0	0	0	0	0	0
	1	-1	-1	-1	0	-1		1	0	0	0	0	0	100
	2	-1	-1	-1	0	-1		2	0	0	0	0	0	0
	3	-1	0	0	-1	0		3	0	80	0	0	0	0
	4	0	-1	-1	0	-1		4	0	0	0	0	0	0
	5	-1	0	-1	-1	0		5	0	0	0	0	0	0

- that
-
- ```
graph TD; 1((1)) --- 3((3)); 1((1)) --- 5((5)); 3((3)) --- 2((2)); 3((3)) --- 4((4)); 2((2)) --- 4((4));
```

- $Q(3,1) = R(3,1) + 0.8 * \max[Q(1,3), Q(1,5)]$   
 $= 0 + 0.8 * \max[0, 100] = 0 + 80 = \mathbf{80}$

- Update the Matrix Q.
- |   | 0  | 1  | 2  | 3  | 4  | 5   |
|---|----|----|----|----|----|-----|
| 0 | -1 | -1 | -1 | -1 | 0  | -1  |
| 1 | -1 | -1 | -1 | 0  | -1 | 100 |
| 2 | -1 | -1 | -1 | 0  | -1 | 100 |
| 3 | -1 | 0  | 0  | -1 | 0  | -1  |
| 4 | 0  | -1 | -1 | 0  | -1 | 100 |
| 5 | -1 | 0  | -1 | -1 | 0  | 100 |

|       | Action |    |   |   |   |     |
|-------|--------|----|---|---|---|-----|
| State | 0      | 1  | 2 | 3 | 4 | 5   |
| 0     | 0      | 0  | 0 | 0 | 0 | 0   |
| 1     | 0      | 0  | 0 | 0 | 0 | 100 |
| 2     | 0      | 0  | 0 | 0 | 0 | 0   |
| 3     | 0      | 80 | 0 | 0 | 0 | 0   |
| 4     | 0      | 0  | 0 | 0 | 0 | 0   |
| 5     | 0      | 0  | 0 | 0 | 0 | 0   |